Development of an Application for Detecting Phishing and Fraudulent Messages

FraudShield
Development of an Application for Detecting Phishing and Fraudulent Messages

First probationary period submission materials
Software Engineering
20 October 2026

Team
Nurassyl Auesbekuly Kendebaev
Askhat Nurlanuly Amirkhanov

Supervisor
Azamat Galymzhanovich Serek
Associate Professor

The project supports manual checking of suspicious text and links, with explanations of risk indicators. This submission presents domain analysis, requirements, architecture, mockups and an evaluation plan.

Status: working draft for team and supervisor review. User interviews, verified test results and substantive supervisor feedback require actual evidence. Section 11 lists the outstanding evidence.

1 Project proposal

Problem and target audience

A user receives a message claiming to come from a bank, delivery company or public service and must decide whether to trust its request, link or demand for urgent action. The proposed system helps assess the message before opening a link. The primary audience is mobile device users in Kazakhstan who read messages in Russian or Kazakh.

Aim

Develop and evaluate an application prototype that accepts message text, identifies phishing and fraud indicators, examines the structure of embedded links and explains the risk in the selected interface language.

Objectives

1. Study smishing detection approaches and limitations of available datasets.
2. Validate user needs and agree measurable requirements.
3. Prepare an anonymised corpus and annotation guidelines.
4. Implement the client, API, opt-in history and analysis module.
5. Compare rules, a simple text classifier and an extended approach.
6. Evaluate detection quality, false alarms, latency and explanation clarity.

Scope

The MVP includes manual text entry, URL extraction without visiting links, explanations and optional check history. The main channel is a mobile web application; a Telegram Mini App is considered as a wrapper. Automatic access to SMS or private chats, call analysis, OCR, payment blocking and guarantees of message safety are outside the MVP.

Engineering challenges and expected outcome

Challenges include combining text and URL analysis, handling two languages and mixed messages, maintaining component contracts, protecting privacy and evaluating reproducibly without data leakage. Expected outputs are an integrated prototype, repository, documentation and a report with measured results. Superiority over existing products is not claimed before testing.

2 Relevance and user needs validation

Official Kazakhstan materials describe phishing, fake websites and messenger fraud as current threats [10]. This establishes the problem, but does not by itself demonstrate demand for this application. Research highlights the importance of local language, context and data [1, 2, 7].

ID	Statement	Evidence and status
E1	Phishing and service impersonation are current problems	Supported by official source [10].
E2	Local data matters for evaluation	Supported by research [1, 2, 7]; applicability to our corpus must be evaluated.
H1	Users will manually paste messages	Hypothesis: validate in discussions and a mockup task.
H2	An explanation is more useful than colour alone	Hypothesis: assess understanding of the result.
H3	Telegram is preferable to a separate app	Hypothesis: ask about the preferred channel.
H4	Most users need check history	Hypothesis: assess usefulness and willingness to store text.

User discussion plan

Before freezing requirements, conduct 3-5 discussions with audience representatives: a student, a working banking-app user, an adult with limited digital experience and, if available, a support specialist. Participation is voluntary. Do not collect passwords, verification codes, financial details or identifiable private correspondence.

Ask about the most recent suspicious message, how it was checked, difficulties, willingness to share data, preferred language and channel, interpretation of "no clear risk indicators detected" and actions after a warning. Finish with an unassisted mockup task.

For each discussion, record the date, anonymised participant code, audience category, consent, actual responses, findings and related requirements. Store records in evidence/. Record requirement changes as decisions. No actual discussion records were supplied; this item remains outstanding.

3 Scientific literature review

The review covers eight scientific works [1-8], mainly from 2024-2026; XLM-R [8] is retained as foundational research. Full texts were examined for problem definition, data, evaluation and limitations. Posters and preprints are explicitly labelled. Published experiments were not reproduced at this stage; implications for FraudShield are design decisions rather than measured results.

Source	Approach and useful findings	Limitation and project implication
[1] Cejas et al., 2026	Multilingual embeddings for HAM, SPAM and SMISHING; local execution with latency and memory evaluation.	Private corpus; no per-language metrics or concept-drift evaluation. Evaluate language and temporal slices.
[2] Ali et al., 2025	Crowdsourced local messages; zero/few-shot LLMs, per-class F1 and macro-F1.	Mozambique and Portuguese; a static time slice. Adapt the collection protocol, not the reported performance, to Kazakhstan.
[3] Timko, Rahman, 2024	SmishTank: user-submitted smishing examples, brands, domains and reputation at collection time.	Contents vary by version; a phishing corpus does not supply legitimate examples. Record snapshot, date and label provenance.
[4] Nahapetyan et al., 2024	Study of tactics and infrastructure; published messages and campaign groupings.	Similar campaigns can leak across splits. Group templates and domains during evaluation.

The corpus must record source, language, date, campaign and label provenance as well as text. Foreign-language studies provide methodological guidance, but do not establish performance on Russian and Kazakh messages.

3 Scientific literature review continued

Source	Approach and useful findings	Limitation and project implication
[5] Hosseinpour, Das, 2025	Four feature streams: regional entities, structure, CharCNN and context; comparison with individual streams.	Short poster; some spam is labelled smishing using keywords, risking dependence on label rules. Use independent annotation and ablations.
[6] Islam et al., 2024	21 extracted URL features and a CNN; explanation of feature contributions.	URL classification omits messages without links. High accuracy on the authors' split does not establish generalisation to new campaigns.
[7] Togmanov et al., 2025	KazMMLU evaluates Kazakh, Russian and regional knowledge, showing model limitations.	Knowledge questions are not phishing detection. Use it to motivate language evaluation, not as fraud training data.
[8] Conneau et al., 2020	XLM-R: multilingual pretraining, transfer and the trade-off between language coverage and model capacity.	General NLP benchmarks do not guarantee fraud detection. Compare quality and resources against TF-IDF on the same split.

Synthesis and research question

Does combined text and URL-structure analysis provide practically useful fraud detection on a local RU/KZ corpus at an acceptable false-positive rate and response time? Compare it with rules and a simple text classifier. Finding no improvement is also a valid research result.

The intended practical focus is local examples, two languages and understandable explanations. These features alone do not establish scientific novelty: multilingual and locally adapted approaches already exist [1, 2, 5].

Additional sources

UCI SMS Spam Collection [9] supports a reproducible baseline, but spam/ham is not equivalent to fraud/legitimate. Product documentation [11-14] supports the system comparison; technology documentation [15-20] supports architecture decisions. Full references appear at the end.

4 Comparison of existing systems

Four systems are compared against eight common criteria using official descriptions [11-14], checked on 20 October 2026. "Not verified" means no supporting evidence was found in the reviewed source; it does not establish that a feature is absent. Products have not yet been tested on a local corpus.

Criterion	Scamio [11]	Norton Genie [12]	VirusTotal [13]	Safe Browsing [14]
Full message text	Yes	Yes	Not the service's primary input	No; URL-focused service
URL	Yes	Yes	Yes	Yes
Image / screenshot	Yes	Yes	Not verified for message explanation	Not applicable to text analysis
User explanation	Analysis and advice described	Yes; explains indicators	Engine and reputation reports	URL threat status
RU/KZ quality on local scams	Not verified	Not verified	Not verified	Does not assess text meaning
API for our application	Not verified	Not verified	API available; verify terms	API available; usage restrictions
Local processing deployment	Not verified	Not verified	Cloud service	Cloud checking
Comparison role	Direct functional alternative	Direct functional alternative	URL-component alternative	URL-component alternative

Rationale for building FraudShield

Scamio and Genie already address similar user needs. The diploma focuses on engineering and evaluating an independently developed system with a controlled corpus, RU/KZ scenarios, an explicit storage policy and reproducible baselines. Any accuracy, language or privacy advantage must be demonstrated.

Before integrating external URL services, verify pricing, terms, data transfers and API restrictions. The MVP can analyse URL structure locally without sending message text to third parties. Absence from a reputation database does not establish safety.

5 Technology selection

These are provisional decisions for team discussion. Finalise the stack after reviewing team skills and a short technical prototype. Planned responsibilities follow the supplied project proposals.

Decision	Alternatives	Proposed choice and rationale
Server API	Node.js + Express; Python + FastAPI; ASP.NET Core	Express for the user API and FastAPI for analysis fit the roles and Python ecosystem. Trade-off: two services. Consolidate in FastAPI if maintenance becomes difficult [15,16,20].
Storage	PostgreSQL; SQLite; MongoDB	PostgreSQL supports user, check, version and feedback relationships; JSONB stores explanations [17]. SQLite is simpler locally; MongoDB suits documents but offers no clear advantage for this relational model.
Client channel	Mobile web / React; Telegram Mini App; native app	Responsive web as the main client; Telegram Mini App as a wrapper [18]. No app installation is required; evaluate Telegram authentication and platform constraints separately.
Text analysis	Rules; TF-IDF + logistic regression; multilingual encoder	Rules as B0, TF-IDF as B1 [19], encoder as candidate B2 [1,8]. Choose B2 based on measurements. An external LLM requires a separate consent and cost experiment.
Component communication	REST/JSON; gRPC; task queue	REST/JSON offers straightforward debugging and contract checks for manual MVP requests. gRPC fits stricter requirements; a queue fits long-running asynchronous jobs.

Reconsideration criteria

Assess integration effort, reproducible builds, latency, cost and the ability to keep private text out of external services. If two server stacks add little value, consolidate the API. Next.js is optional: SEO and server rendering are not primary requirements for the checking form.

6 Functional requirements

Must means required for the MVP; Should indicates the next priority. Requirements remain provisional until user discussions and supervisor approval. Acceptance criteria describe future checks, not completed tests.

ID	Priority	Requirement	Acceptance criterion
FR1	Must	RU/KZ text input, 1-5000 characters	Reject empty and over-limit input; accept valid Unicode strings.
FR2	Must	Extract and parse links	Extract all URLs in the control set without automatically visiting or downloading sites.
FR3	Must	Analysis result	Return risk_level, reasons, model_version and status. A failure returns status=unavailable, without a green result.
FR4	Must	Understandable explanation	Each detected signal has an explanation and advice; absence of signals is not a safety guarantee.
FR5	Must	RU/KZ interface selection	Display input controls, results, errors and guidance in the selected language; do not silently translate the original text.
FR6	Must	History with explicit consent	Do not store text by default. Consented records are available only to their owner.
FR7	Must	Delete a saved check	The owner can delete a record; subsequent reads return 404; prohibit access by other users.
FR8	Should	Result feedback	Keep feedback separate from expert labels; do not automatically add it to training data.
FR9	Must	Analysis version tracking	Saved results include algorithm version, timestamp and response schema version.
FR10	Should	Telegram wrapper	If enabled, verify initData signatures server-side; reject forged or expired data.

6 Non-functional requirements

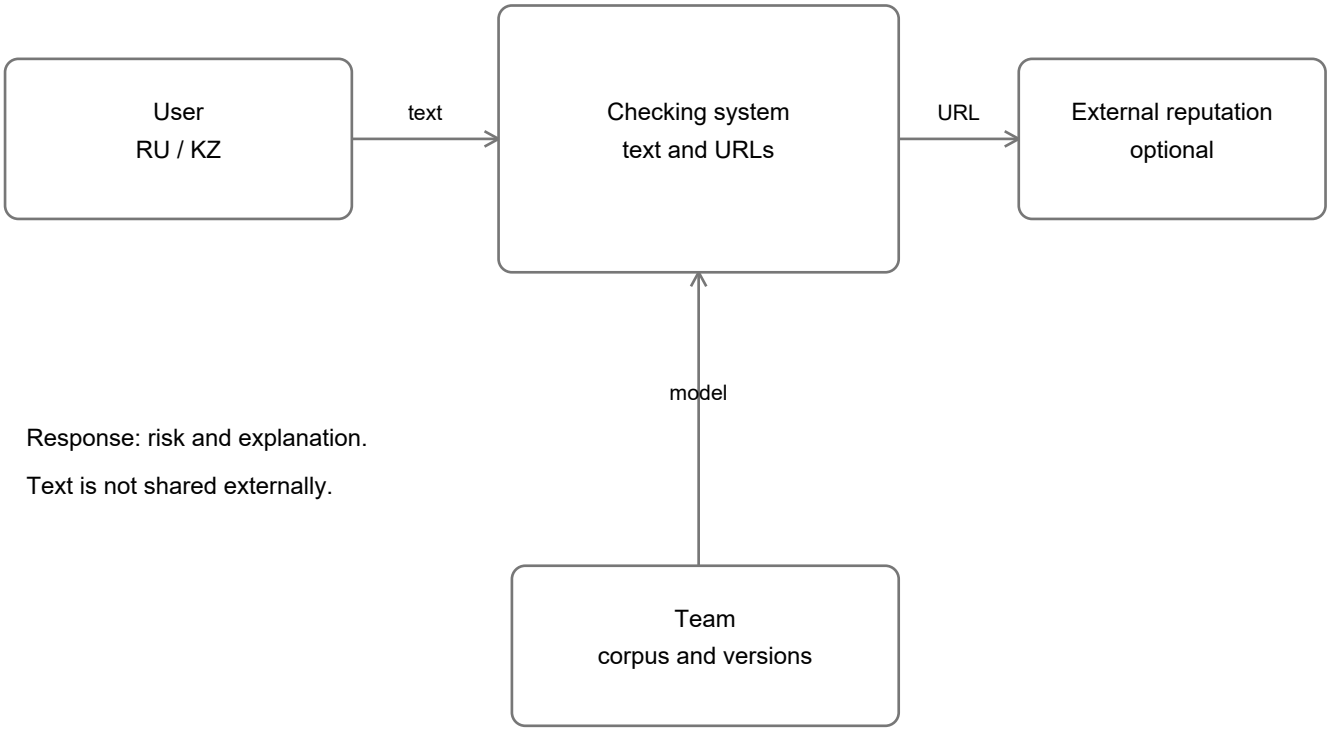
ID	Priority	Measure	Conditions and acceptance
NFR1	Must	Local analysis p95 ≤ 2 s	1000 requests after warm-up, 10 concurrent clients, 100-2000 characters; record CPU/RAM and versions. Not measured yet.
NFR2	Must	Timeout ≤ 5 s	After timeout, display unavailability and retry; never substitute a safe result.
NFR3	Must	Data minimisation	Do not store raw text without consent or write it to logs. Inspect logs and database using control strings.
NFR4	Must	User isolation	Reject API reads/deletions of another user's history; keep secrets outside the repository.
NFR5	Must	Reproducibility	Clean installation from lock files and README test commands; record model versions and split manifest.
NFR6	Must	Quality by language	Report precision, recall, F1 and confusion matrices separately for RU and KZ; identify mixed-language messages.
NFR7	Should	Interface clarity	4 of 5 participants complete a check unassisted; report actual counts and comments without population-wide generalisation.

Result wording policy

Internal categories are high, suspicious and low_signal. User-facing labels are "High risk", "Suspicious indicators" and "No clear risk indicators detected". Service failures have a separate state. The original "safe" label is replaced because classification errors are unavoidable. Do not display numerical probabilities until calibration has been validated.

7 Architecture and use cases

System context

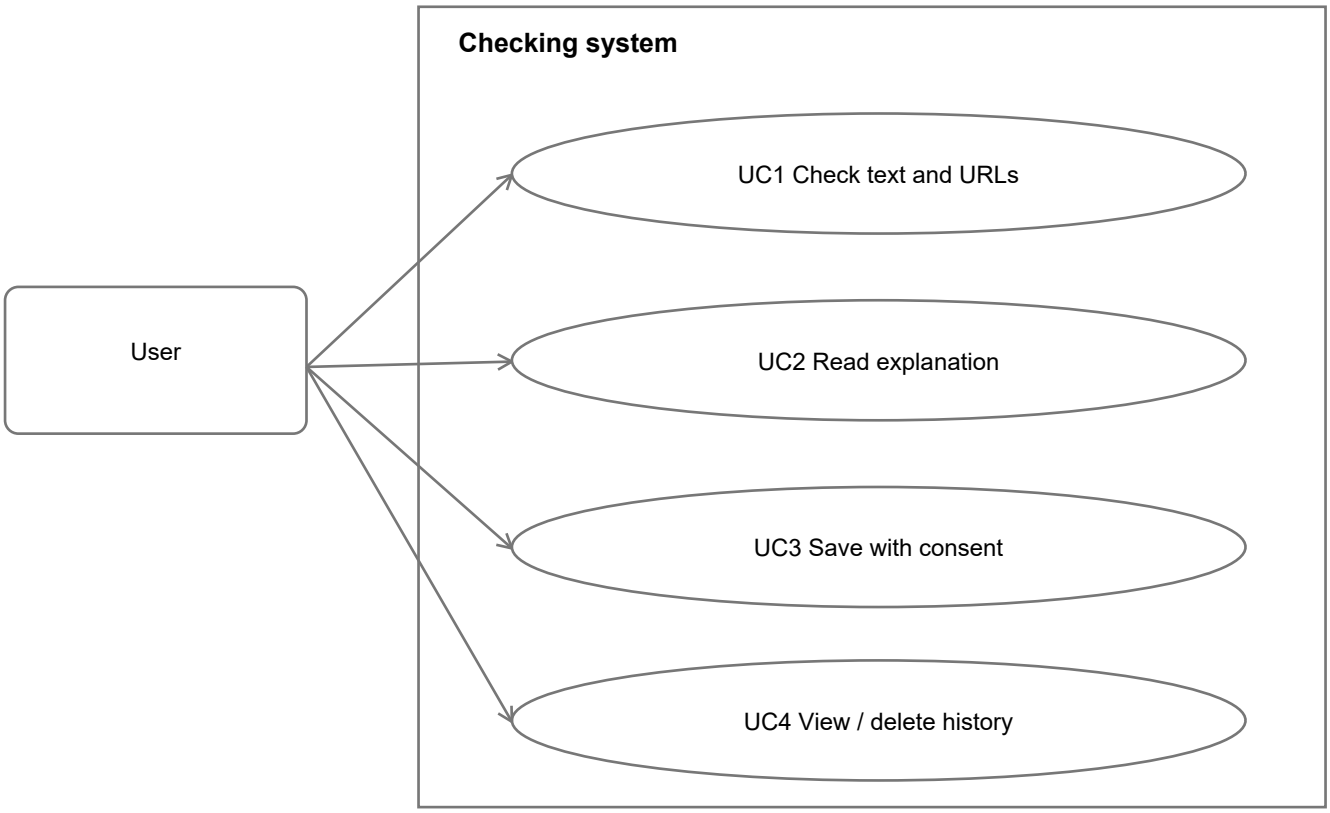


Users submit text manually and receive an assessment. The team manages corpus and model versions. External URL reputation checks are optional; the default analysis does not visit links. The architecture requires no access to private chats or SMS.

Main flow UC1

The user pastes a message → the interface checks length → the API validates the request → the analyser normalises text and extracts URLs → signals are combined → an explanation is returned → the interface displays the result. Saving requires consent and an authenticated owner. A timeout produces an availability error.

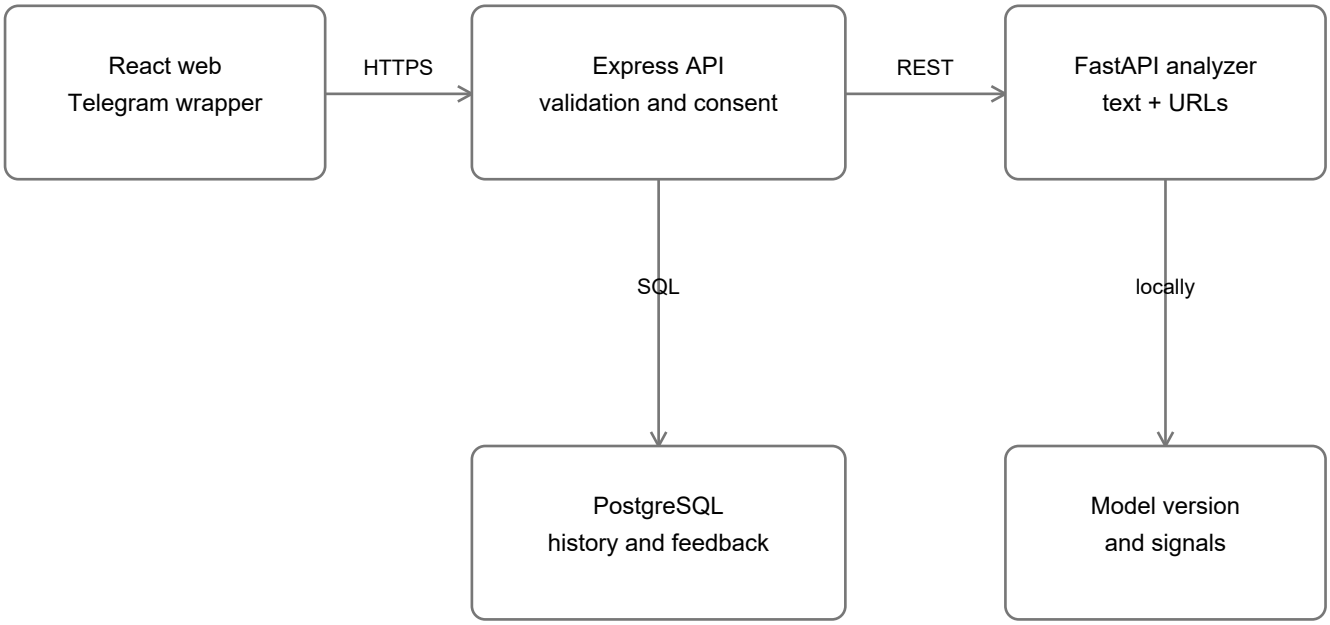
7 Use case diagram



Use case	Precondition	Success / exception
UC1 Check	Text within length limits	FR1-3. A result or explicit error without a safety assessment.
UC2 Explanation	Completed analysis	FR4-5. The user sees indicators and next steps.
UC3 Save	Sign-in and explicit consent	FR6, FR9. The record belongs to the current user.
UC4 History and deletion	Signed-in user	FR7. Other users' records are inaccessible; deleted records cannot be read.

Feedback FR8 is an optional action after a result. Telegram authentication FR10 applies only when the wrapper is enabled. Anonymous checks do not depend on Telegram.

7 Components and data



Proposed architecture. Service implementation and integration are not yet verified.

Entity	Purpose and main fields
User	id, auth_provider, provider_subject; avoid unnecessary profile data.
Check	id, user_id, created_at, risk_level, reasons_json, model_version, retention_until; store text only with separate consent.
Feedback	id, check_id, user_id, verdict, comment; keep user opinions separate from expert labels.
DatasetItem	id, sanitized_text, language, label, source, source_date, consent_or_license, campaign_group, split, annotation_version.

Anonymous checks do not create database history. A retention period must be explicit; the provisional proposal is 30 days for opt-in history with immediate deletion available. Agree this period with the team and supervisor before implementation.

8 Contracts and traceability

Draft API contract

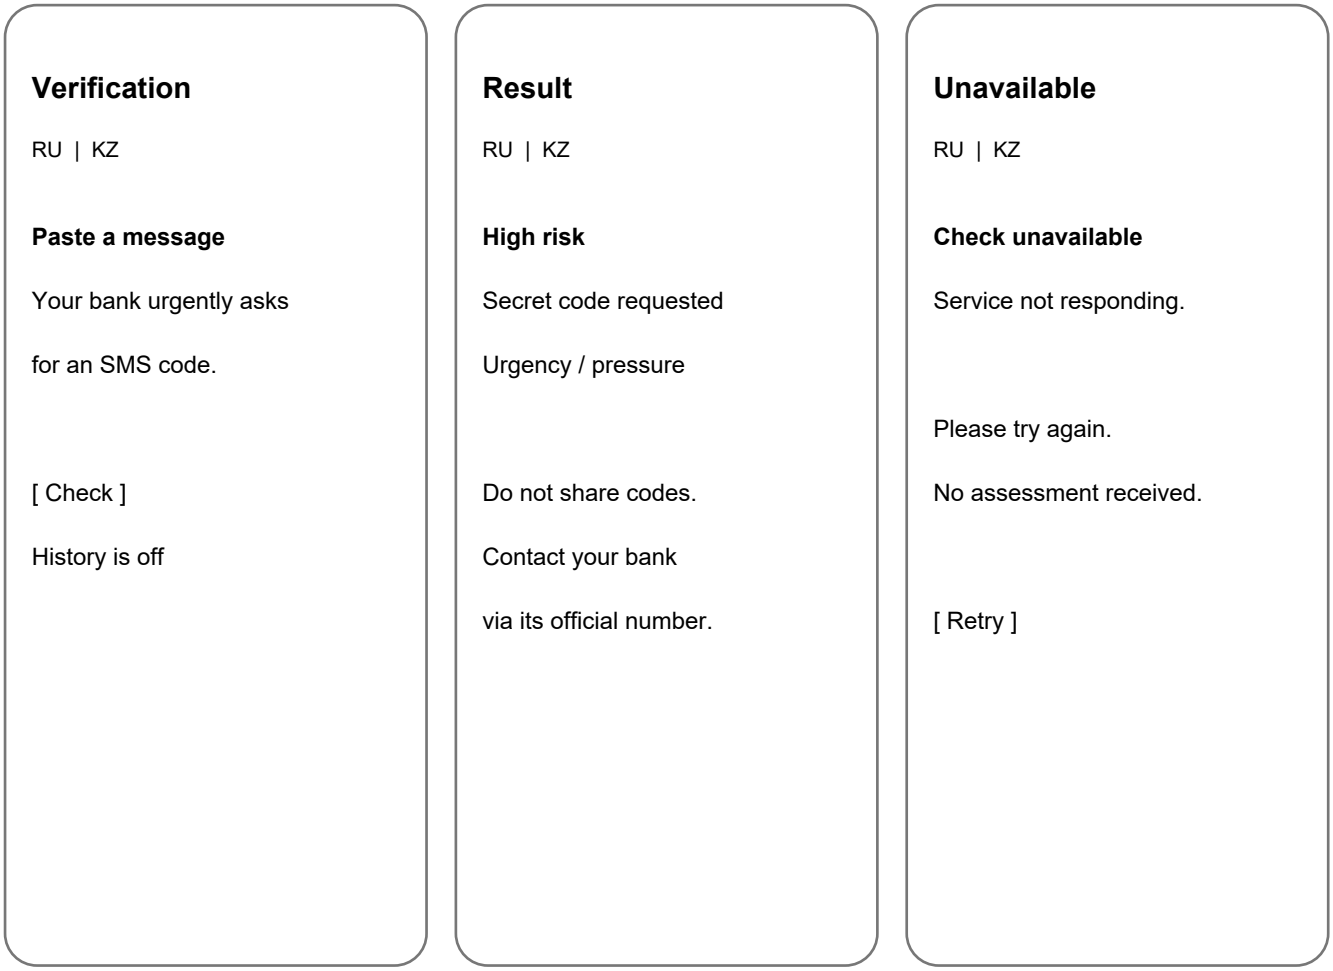
POST /api/v1/checks/analyze
Request: text (1-5000 characters), language (ru/kk/auto). Anonymous message text is not stored.
Response 200: request_id, status=completed, risk_level, reasons[], urls[], model_version, processing_ms. Reasons contain code and description.
Errors: 422 invalid input; 429 rate limit exceeded; 503 analysis unavailable. A failure must not return risk_level=low_signal.

POST /api/v1/history: save a result with explicit consent for an authenticated user.
GET /api/v1/history: current owner's history.
DELETE /api/v1/history/{id}: delete an owned record.
POST /api/v1/feedback: optional feedback; does not automatically change the reference label.

Requirements, components and verification

Requirement	Component	Verification
FR1, FR5	Web + API	T1: empty input, length boundaries, Unicode, language switching.
FR2	URL analyzer	T2: scheme-less URLs, IDN, IP, userinfo and malformed addresses; no network requests.
FR3, FR4, FR9	Analyzer + UI	T3: response schema, version, reasons and service failure.
FR6, FR7, NFR3-4	Auth + API + PostgreSQL	T4: consent, another user's record, deletion and no text in logs.
NFR1-2	API + Analyzer	T5: load and timeout on documented hardware.
NFR6	Data pipeline + evaluator	T6: fixed split; metrics by RU/KZ and class.
NFR7	Web prototype	T7: tasks with 5 users and an error record.

9 UX and mockups



Design mockups. Demonstration text is not user data.

Key interactions

The first screen explains the purpose, input limit and absence of automatic storage. Disable checking for empty input and while a request is running. Results contain a text category, specific indicators and practical advice; colour is only a supplementary cue.

The low-signal state retains: "This does not guarantee safety. Verify the sender through an official channel." Error screens do not show a green checkmark. History controls appear after sign-in; consent to history storage does not imply consent to research use.

Mockups cover UC1-2 and the error state. History, Kazakh localisation and Telegram sign-in require further design. Detailed interaction specifications are in prototype/wireframes.md. User validation has not yet taken place.

10 Verification and evaluation protocol

Data and annotation

Planned pilot: 600 anonymised messages, 300 each for RU and KZ, with both classes represented. Initial target per language: 100 fraudulent and 100 legitimate messages. These are collection targets, not an existing dataset. If data is insufficient, report actual coverage and uncertainty instead of replacing authentic data with generated examples.

Sources: lawfully available public corpora; voluntary anonymised examples; controlled artificial examples for unit tests, labelled as synthetic. Remove identifiers and secret codes. Retain language, date, campaign and URL signals; do not automatically open malicious links.

Use two reference classes: fraud and legitimate; uncertain cases require review. Two annotators independently label the first 100 examples, discuss disagreements and refine guidelines. Keywords and suspicious URLs are evidence for annotation, not automatic proof of the true class. Keep advertising spam separate from fraud.

Splits and baselines

After deduplication, create approximately 60/20/20 train/validation/test splits grouped by campaigns and domains, with language balance. Copies of the same template or campaign must not cross splits. Record the seed, identifiers and configuration hash. Select thresholds on validation only. If dates permit, additionally evaluate a later time slice.

B0: deterministic text and URL rules.
B1: word/character TF-IDF with logistic regression.
B2: multilingual encoder with URL signals. Evaluate after B0/B1; compare text-only, URL-only and combined variants on the same test set.

10 Metrics and acceptance criteria

Verification	Metric and procedure	Provisional criterion
Detection	Fraud precision, recall and F1; macro-F1; confusion matrix, sample sizes and 95% CI; separately for RU/KZ.	Target recall ≥ 0.85 and FPR ≤ 0.10 per language. Agree thresholds; no results measured yet.
Uncertainty	Share of suspicious responses and coverage; compute FPR/recall across binary decisions and report the threshold.	Do not conceal difficult cases by excluding them; report abstention separately.
Latency	p50/p95/p99 and error rate; 1000 requests, 10 clients, local analysis, fixed CPU/RAM and versions.	NFR1: p95 ≤ 2 s; timeout ≤ 5 s.
Privacy and API	History isolation, deletion, no raw text in logs; HTML and instruction-injection attempts.	Mandatory T1-T4 checks pass; the analyser does not open links.
User flow	5 volunteers; unassisted completion, task time and warning comprehension.	4/5 complete a check; disclose all errors and small-sample limitations.
Comparison	B0/B1/B2 on the same test split; paired bootstrap for the F1 difference; document configurations.	Report actual improvement or its absence; do not import results from other studies.

Reporting results

Report configurations, dataset sizes and composition, threshold selection, metrics and limitations. No local quality measurements have been performed yet. If 600 examples do not provide sufficient statistical power, report wide confidence intervals and a plan to expand the corpus.

11 Project management and risks

Participant	Planned responsibility	Verified contribution at this stage
Nurassyl Auesbekuly Kendebaev	Client, API, PostgreSQL, initial data collection and integration.	Initial repository and working documentation prepared with AI assistance. Individual verification and team review are pending.
Askhat Nurlanuly Amirkhanov	Text classification, URLs and domains, experiments and metrics.	The supplied proposal states this role. Commit history and actual completed tasks have not yet been verified.

Repository: https://github.com/Kendebaev/FraudShield
YouTrack: https://frauddetect.youtrack.cloud/issues/DIP
Accounts: Kendebaev and askmirkhanov-boop.

Working process

One repository serves both team members. Each task has an owner, expected outcome, acceptance criteria and deadline. Include its ID in the branch and commit/PR description; provide genuine commits under the contributor's account. Changes undergo a second-member review. Main contains agreed versions; intermediate drafts stay on branches.

The supervisor and commission need access to all materials. The current repository is publicly readable; verify reviewer access to private YouTrack resources or explicitly agree another accessible evidence route. Confirm access before submitting the form.

Immediate plan

21 October: discussions, evidence and scope.
22 October: requirements traceability, data and architecture.
23 October: supervisor review, corrections and access checks.
24 October: final version, prob-period-1 tag, board PDF and form before 18:00. Recommended internal readiness deadline: 16:00 in the deadline's applicable time zone.

Tasks beyond 24 October are provisional and do not replace the immediate plan. Fifteen tasks with criteria exist in YouTrack; management/tasks.csv is their snapshot. Infrastructure tasks are in progress, prepared documents are in review and remaining items are in the backlog. Completion is not claimed. The second member has no separate assigned YouTrack account yet; responsibility is recorded in task descriptions.

11 Risks and submission readiness

Risk	Owner	Mitigation and review trigger
Insufficient real Kazakh data	A. N. Amirkhanov	Start a pilot early; evaluate RU/KZ separately. Limit claims and expand collection when data is insufficient.
Template leakage across splits	A. N. Amirkhanov	Deduplicate and group campaigns/domains; manually inspect similar messages.
False sense of safety	Both members	Use cautious wording, a separate error state and measurement of missed fraud.
Disclosure of private text	N. A. Kendebaev	No storage by default; separate consent; no automatic external text analysis.
Complexity of two server stacks	N. A. Kendebaev	Define contracts and integration smoke checks; consolidate the API if delays arise.
Inaccessible submission links	N. A. Kendebaev	Verify with reviewer access; confirm permissions before submission.

Outstanding evidence

1. Each member's reviewed contribution and personal YouTrack access.
2. Actual records of 3-5 discussions or documented client agreement.
3. Substantive supervisor feedback and linked responses.
4. Genuine individual contributions and team review.
5. Updated PDF screenshots of the final board.
6. Reviewer access to repository and tracker, a final tagged version and the submission form.

The proposals provided are not evidence of user discussions, experiments or code contributions. The form and screenshot PDF are prepared separately; access and formal submission remain unconfirmed.

12 Defence preparation

Suggested project introduction

We are developing an assistant for checking suspicious messages in Russian and Kazakh. Users paste text manually and receive risk indicators, explanations and advice. The diploma will evaluate rules, a simple text classifier and a multilingual model with URL signals. One member covers the client, API and database; the other covers analysis and evaluation. Quality will be measured on a fixed local corpus against reproducible baselines.

Repository structure and submission package

Question	Evidence for the answer
Who needs this project and why?	Problem evidence [10] and actual discussions once conducted. Distinguish hypotheses from findings.
How does it differ from Scamio and Genie?	Independent engineering and verifiable local evaluation; superiority is not yet demonstrated.
Why this stack?	Section 5, requirements and integration costs. Simplify based on prototype findings.
What does suspicious mean?	An uncertainty region and indicators, not a guaranteed diagnosis or an unqualified third ground-truth class.
How will you establish quality?	B0/B1/B2, grouped test split, RU/KZ metrics, FPR, recall, confidence intervals and latency.
What did each member do?	Show personal tasks, changes and reviews; explain every decision.
What already works?	Demonstrate only code that actually runs. These materials cover analysis and design, not a complete application.

The repository contains a README, report, diagrams, task snapshot and editable sources for team review. Each participant should commit genuine work under their own account. The submission requires two PDFs: the unified report and real board/task screenshots.

References

[1] Cejas O. A., Guo Y., Tang Q. (2026). **Large-Scale Multilingual SMS Fraud Detection For Telecom Networks**. NLPAICS, scientific publication.
https://aclanthology.org/2026.nlpaics-1.7/

[2] Ali F. D. M. A. et al. (2025). **MOZ-Smishing: A Benchmark Dataset for Detecting Mobile Money Frauds**. AfricaNLP, scientific publication.
https://aclanthology.org/2025.africanlp-1.23/

[3] Timko D., Rahman M. L. (2024). **Smishing Dataset I: Phishing SMS Dataset from Smishtank.com**. CODASPY, scientific publication; author version on arXiv.
https://arxiv.org/abs/2402.18430

[4] Nahapetyan A. et al. (2024). **On SMS Phishing Tactics and Infrastructure**. IEEE S&P, scientific publication; authors' materials.
https://github.com/wspr-ncsu/sms-phishing

[5] Hosseinpour S., Das S. (2025). **POSTER: A Multi-Signal Model for Detecting Evasive Smishing**. WiSec, scientific poster; author version on arXiv.
https://arxiv.org/abs/2505.18233

[6] Islam M. R. et al. (2024). **PhishGuard: A Convolutional Neural Network Based Model for Detecting Phishing URLs with Explainability Analysis**. Scientific preprint on arXiv.
https://arxiv.org/abs/2404.17960

[7] Togmanov M. et al. (2025). **KazMMLU: Evaluating Language Models on Kazakh, Russian, and Regional Knowledge of Kazakhstan**. ACL, scientific publication.
https://aclanthology.org/2025.acl-long.701/

[8] Conneau A. et al. (2020). **Unsupervised Cross-lingual Representation Learning at Scale**. ACL, foundational scientific publication.
https://aclanthology.org/2020.acl-main.747/

[9] UCI Machine Learning Repository (2012). **SMS Spam Collection**. Dataset.
https://archive.ics.uci.edu/dataset/228/sms+spam+collection

[10] Government portal of the Republic of Kazakhstan (2026). **Law and order and protection of citizens' digital rights [translated title]**. Official information on fraud.
https://www.gov.kz/org/aqmola-selinograd/press/news/details/1308837?lang=ru

Access and documentation check date: 20 October 2026. The first eight entries [1-8] are scientific works; the others are data, official materials and documentation. Links continue on the next page.

References continued

[11] Bitdefender (n.d.). **Scamio**. Official product description.
https://www.bitdefender.com/en-sg/consumer/scamio

[12] Norton (n.d.). **FAQ: Norton Genie**. Official product support.
https://support.norton.com/sp/en/us/home/current/solutions/v20230717145233467

[13] VirusTotal (n.d.). **How it works; Private Scanning**. Official documentation.
https://docs.virustotal.com/docs/how-it-works

[14] Google (n.d.). **Safe Browsing**. Official documentation.
https://developers.google.com/safe-browsing

[15] FastAPI (n.d.). **Features**. Official documentation.
https://fastapi.tiangolo.com/features/

[16] Express (n.d.). **Express web framework**. Official documentation.
https://expressjs.com/

[17] PostgreSQL (n.d.). **JSON Types**. Official documentation.
https://www.postgresql.org/docs/current/datatype-json.html

[18] Telegram (n.d.). **Telegram Mini Apps**. Official documentation.
https://core.telegram.org/bots/webapps

[19] scikit-learn (n.d.). **Text feature extraction**. Official documentation.
https://scikit-learn.org/stable/modules/feature_extraction.html#text-feature-extraction

[20] Microsoft (n.d.). **ASP.NET Core APIs overview**. Official documentation.
https://learn.microsoft.com/en-us/aspnet/core/fundamentals/apis?view=aspnetcore-10.0

Additional VirusTotal privacy documentation: https://docs.virustotal.com/docs/private-scanning. Published product descriptions require practical verification; pricing and usage conditions have not been assumed.